

# Psy i koty

... czyli klasyfikacja obrazów i manipulacja dużą ilością danych

## Karolina Kotlowska

### 13.1. Pobieramy dane

Załadujemy zbiór danych z kolekcji [TensorFlow Datasets](https://www.tensorflow.org/datasets) (<https://www.tensorflow.org/datasets>). Zbiór jest widoczny jako obiekt klasy `tf.data.Dataset` zapewniającej funkcjonalność podobną do strumieni.

Przy pierwszym uruchomieniu dane zostaną pobrane. Później będzie można je czytać wielokrotnie.

In [5]:

```
import matplotlib.pyplot as plt
import numpy as np
import tensorflow as tf
import tensorflow_datasets as tfds

# setattr(tfds.image_classification.cats_vs_dogs, '_URL', "https://download.microsoft.com/download/3/E/1/3E1C3F21-ECDB-4869-8368-6DEBA77B919F/kagglecatsanddogs_5340.zip")
ds = tfds.load('cats_vs_dogs', split='train', batch_size=None, as_supervised=True)
```

Liczba obrazów to:

In [6]:

```
len(ds)
```

Out[6]:

23262

#### 13.1.1 Wyświetlmy przykładowe obrazy

Przed wyświetleniem zostaną przeskalowane do tych samych rozmiarów

In [7]:

```
plt.rcParams["figure.figsize"] = (10,10)
it = ds.as_numpy_iterator()
labels=['Cat', 'Dog']
input_size=100
for i in range(9):
    ax = plt.subplot(330 + 1 + i)
    image, label = next(it)
    image = tf.image.resize(image, size=[256,256], method=tf.image.ResizeMethod.NEAREST_NEIGHBOR)
    ax.set_title(labels[label])
    plt.imshow(image)
plt.show()
```



### 13.1.2 Załadujemy obrazy i zapiszemy w tablicach numpy

Obraz jest trójwymiarową tablicą (szerokość, wysokość, kanały koloru). Aby umieścić wszystkie obrazy w jednej tablicy 4D muszą mieć identyczne rozmiary.

In [8]:

```
import matplotlib.pyplot as plt
import numpy as np
import tensorflow as tf
import tensorflow_datasets as tfds
# from keras.preprocessing.image import img_to_array
from tensorflow.keras.utils import img_to_array
from tqdm import tqdm

def load_images_to_numpy(name,resize_to=[80,80]):
    ds = tfds.load(name, split='train', batch_size=None,as_supervised=True)
    images = []
    labels=[]
    for image, label in tqdm(ds):
        image = tf.image.resize(image, size=resize_to, method=tf.image.ResizeMethod.
NEAREST_NEIGHBOR)
        image = img_to_array(image)/255
        images.append(image)
        labels.append(label.numpy())

    # return images,labels
    X=np.array(images)
    y=np.array(labels)
    return X,y

X,y = load_images_to_numpy('cats_vs_dogs')


print(X.shape)
print(y.shape)
```

100%|██████████| 23262/23262 [00:58&lt;00:00, 400.21it/s]

```
(23262, 80, 80, 3)
(23262,)
```

In [9]:

```
plt.imshow(X[0]);
```



## 13.2. Klasyfikacja zbioru załadowanego do pamięci

Wypróbujemy kilka klasyfiaktorów z biblioteki sklearn. Przedtem musimy spłaszczyć X do postaci dwuwymiarowej tablicy.

Nie oczekujemy tu dobrych wyników, ponieważ standardowe algorytmy traktują obraz, jak długi wektor pikseli, a nie biorą pod uwagę cech lokalnych.

Gdyby przeskalować obraz do rozmiarów 4x4 otrzymalibyśmy informacje o kolorach w różnych regionach obrazu. Może pozwoliłoby to odróżnić krajobraz morski i lądowy, ale nie psy i koty.

In [10]:

```
X = np.reshape(X, (X.shape[0], -1))  
print(X.shape)
```

```
(23262, 19200)
```

**TODO 13.2.1**

- Wypróbuj wpierw GaussianNB, a potem co najmniej 3 wybrane kolejne klasyfikatory
- Jaka jest efektywność klasyfikacji?
- Co się dzieje z pamięcią
- Zestaw informacje w postaci tabeli

In [11]:

```

from sklearn.model_selection import train_test_split
from sklearn.naive_bayes import GaussianNB
from sklearn.naive_bayes import MultinomialNB
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.linear_model import SGDClassifier
from sklearn.svm import SVC
import sklearn

clf = GaussianNB()
# clf = RandomForestClassifier(n_estimators=100,verbose=1)
# clf = MultinomialNB()
# clf = LogisticRegression()
# clf = KNeighborsClassifier(n_neighbors=5) #umieszcza wszystkie dane w pamięci
# zużywa bardzo pamięć
# clf = SGDClassifier(loss='log')
# clf = SVC(kernel='rbf',gamma='auto')

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.1, random_state = 123)
clf.fit(X_train,y_train)
y_pred = clf.predict(X_test)
print(f'Accuracy:{sklearn.metrics.accuracy_score(y_test,y_pred)}')
print(f'F1:{sklearn.metrics.f1_score(y_test,y_pred,average="macro")}')

```

Accuracy:0.5612376450365277

F1:0.5605205089886541

In [12]:

```
clf = RandomForestClassifier(n_estimators=100,verbose=1)

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.1, random_state = 123)
clf.fit(X_train,y_train)
y_pred = clf.predict(X_test)
print(f'Accuracy:{sklearn.metrics.accuracy_score(y_test,y_pred)}')
print(f'F1:{sklearn.metrics.f1_score(y_test,y_pred,average="macro")}')

[Parallel(n_jobs=1)]: Using backend SequentialBackend with 1 concurrent workers.
[Parallel(n_jobs=1)]: Done 100 out of 100 | elapsed: 4.5min finished
[Parallel(n_jobs=1)]: Using backend SequentialBackend with 1 concurrent workers.

Accuracy:0.6403094112591319
F1:0.6396961880884169

[Parallel(n_jobs=1)]: Done 100 out of 100 | elapsed: 0.3s finished
```

In [13]:

```
clf = LogisticRegression()

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.1, random_state = 123)
clf.fit(X_train,y_train)
y_pred = clf.predict(X_test)
print(f'Accuracy:{sklearn.metrics.accuracy_score(y_test,y_pred)}')
print(f'F1:{sklearn.metrics.f1_score(y_test,y_pred,average="macro")}')

c:\Users\krzyc\anaconda3-1\lib\site-packages\sklearn\linear_model\_logistic.py:814: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:
    https://scikit-learn.org/stable/modules/preprocessing.html
Please also refer to the documentation for alternative solver options:
    https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
    n_iter_i = _check_optimize_result(

Accuracy:0.5762784701332188
F1:0.5762608630130106
```

In [14]:

```

clf = KNeighborsClassifier(n_neighbors=5)

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.1, random_state = 123)
clf.fit(X_train, y_train)
y_pred = clf.predict(X_test)
print(f'Accuracy: {sklearn.metrics.accuracy_score(y_test, y_pred)}')
print(f'F1: {sklearn.metrics.f1_score(y_test, y_pred, average="macro")}')

```

Accuracy: 0.5487752470992694

F1: 0.5253455952630719

Najlepiej sprawdził się RandomForestClassifier.

Trenowanie pochłonęło sporo pamięci. KNeighborsClassifier umieszcza wszystkie dane w pamięci - zużywa bardzo pamięć.

| Model                  | Accuracy  | F1        |
|------------------------|-----------|-----------|
| GaussianNB             | 0.5612376 | 0.5605205 |
| RandomForestClassifier | 0.6438633 | 0.6330582 |
| LogisticRegression     | 0.5758487 | 0.5758108 |
| KNeighborsClassifier   | 0.5487752 | 0.5253456 |

## 13.3. Partial fit

Zbiór wszystkich obrazów załadowany do pamięci zużywa około 30% pamięci środowiska Colab.

Wydzielenie `x_train` i `x_test` to kolejne 30%. W przypadku niektórych klasyfikatorów brak jest miejsca na model...

W przypadku dużych zbiorów danych rozwiązanie może polegać na wykonaniu pętli, w której

- odczytuje się fragment danych (batch) ze źródła
- dane te są użyte do uczenia (uaktualnienia modelu)

W bibliotece `sklearn` wybrane klasyfikatory i algorytmy regresji mają służącą do tego metodę `partial_fit()`.

### 13.3.1 Organizacja odczytu danych

- Obrazy będą czytane z dysku
- Podczas odczytu dokonywana będzie konwersja za pomocą funkcji `img_convert`

In [15]:

```
import matplotlib.pyplot as plt
import numpy as np
import tensorflow as tf
import tensorflow_datasets as tfds
# from keras.preprocessing.image import img_to_array
from tensorflow.keras.utils import img_to_array

from tqdm import tqdm

dataset = tfds.load('cats_vs_dogs', split='train', batch_size=None, as_supervised=True)

def img_convert(image, label):
    image = tf.image.resize(image, size=[80,80], method=tf.image.ResizeMethod.NEAREST_NEIGHBOR)
    image = image / 255
    return (image, label)

dataset = dataset.map(img_convert)
dataset = dataset.shuffle(buffer_size=100)
```

### 13.3.2 Podział na strumień danych uczących i testowych

Jest to podział sztywny. Funkcje `take()` i `skip()` przesuwają granice iteracji, ale nie zapewniają losowego wyboru.

Strumień (dataset) może losowo zmieniać kolejność danych, ale stosując bufor wewnętrzny o pewnym zadanym rozmiarze.

In [16]:

```
train_size = int(0.7*len(dataset))

ds_train = dataset.take(train_size)
ds_test = dataset.skip(train_size)
print(len(ds_train), len(ds_test))

ds_test = ds_test.batch(256, drop_remainder=False)
ds_train = ds_train.batch(128, drop_remainder=False)
```

16283 6979

### 13.3.3 Iteracja po danych

In [17]:

```
for batch, labels in ds_train:
    X=batch.numpy()
```



### 13.3.4 Opcjonalnie: zawartość dataset może zostać zwielokrotniona

Można wówczas wielokrotnie iterować przez te same elementy. Wielokrotna iteracja ma sens, dla modeli wyznaczanych w procesie optymalizacji (np. za pomocą metod gradientowych.) Wówczas przy równoczesnym zastosowaniu losowania możliwe jest opuszczenie minimum lokalnego funkcji kosztu.

In [18]:

```
epochs=2  
ds_train=ds_train.repeat(epochs)
```

### 13.3.5 Klasyfikacja

#### TODO 13.3.1

- Proszę wyznaczyć  $X$  na podstawie danych w zmiennej `batch` oraz  $y$  na podstawie `labels`.  $X$  należy spłaszczyć do tablicy 2-wymiarowej.
- Dla których (poniższych klasyfikatorów) wielokrotne przetwarzanie ('ogładanie') tych samych danych może mieć sens?
- Wypróbuj i porównaj **dwa klasyfikatory** implementujące metodę `partial_fit()`. Jaka jest ich efektywność?

In [19]:

```
from sklearn.naive_bayes import MultinomialNB
from sklearn.naive_bayes import GaussianNB
from sklearn.linear_model import SGDClassifier
from sklearn.linear_model import PassiveAggressiveClassifier

from tqdm import tqdm

import sklearn

# clf = MultinomialNB()
# clf = GaussianNB()
# clf = PassiveAggressiveClassifier()
clf = SGDClassifier(loss='perceptron')

# for i,(batch,labels) in tqdm( enumerate(ds_train) ):
for i,(batch,labels) in enumerate(ds_train) :
    X=batch.numpy()
    X = np.reshape(X,(X.shape[0],-1))
    y=labels.numpy()
    clf.partial_fit(X,y,classes=[0,1])

# predict

y_true_list = []
y_pred_list = []
for i, (batch, labels) in enumerate(ds_test):
    X = batch.numpy()
    X = np.reshape(X, (X.shape[0], -1))
    y = labels.numpy()
    y_pred = clf.predict(X)
    y_true_list.extend(y.tolist())
    y_pred_list.extend(y_pred.tolist())

y_true = np.hstack(y_true_list)
# y_true.shape

y_pred=np.hstack(y_pred_list)
# y_pred.shape

print(f'Accuracy:{sklearn.metrics.accuracy_score(y_true,y_pred)}')
print(f'F1:{sklearn.metrics.f1_score(y_true,y_pred,average="macro")}')
print(f'auroc:{sklearn.metrics.roc_auc_score(y_true,y_pred,average="macro")}')
```

Accuracy:0.5929216220088838

F1:0.5900312108139593

auroc:0.5932807726421603

In [20]:

```
clf = MultinomialNB()

# for i,(batch,labels) in tqdm( enumerate(ds_train) ):
for i,(batch,labels) in enumerate(ds_train) :
    X=batch.numpy()
    X = np.reshape(X,(X.shape[0],-1))
    y=labels.numpy()
    clf.partial_fit(X,y,classes=[0,1])

# predict

y_true_list = []
y_pred_list = []
for i, (batch, labels) in enumerate(ds_test):
    X = batch.numpy()
    X = np.reshape(X, (X.shape[0], -1))
    y = labels.numpy()
    y_pred = clf.predict(X)
    y_true_list.extend(y.tolist())
    y_pred_list.extend(y_pred.tolist())

y_true = np.hstack(y_true_list)
# y_true.shape

y_pred=np.hstack(y_pred_list)
# y_pred.shape

print(f'Accuracy:{sklearn.metrics.accuracy_score(y_true,y_pred)}')
print(f'F1:{sklearn.metrics.f1_score(y_true,y_pred,average="macro")}')
print(f'auroc:{sklearn.metrics.roc_auc_score(y_true,y_pred,average="macro")}')
```

Accuracy:0.5601088981229403

F1:0.5600167487263904

auroc:0.5600504061002419

In [21]:

```

clf = PassiveAggressiveClassifier()

# for i,(batch,labels) in tqdm( enumerate(ds_train) ):
for i,(batch,labels) in enumerate(ds_train) :
    X=batch.numpy()
    X = np.reshape(X,(X.shape[0],-1))
    y=labels.numpy()
    clf.partial_fit(X,y,classes=[0,1])

# predict

y_true_list = []
y_pred_list = []
for i, (batch, labels) in enumerate(ds_test):
    X = batch.numpy()
    X = np.reshape(X, (X.shape[0], -1))
    y = labels.numpy()
    y_pred = clf.predict(X)
    y_true_list.extend(y.tolist())
    y_pred_list.extend(y_pred.tolist())

y_true = np.hstack(y_true_list)
# y_true.shape

y_pred=np.hstack(y_pred_list)
# y_pred.shape

print(f'Accuracy:{sklearn.metrics.accuracy_score(y_true,y_pred)}')
print(f'F1:{sklearn.metrics.f1_score(y_true,y_pred,average="macro")}')
print(f'auroc:{sklearn.metrics.roc_auc_score(y_true,y_pred,average="macro")}')

```

Accuracy:0.5406218655967904

F1:0.46612554829406794

auroc:0.5420751471613843

Porównanie klasyfikatorów:

| Model                       | Accuracy | F1 Score | AUROC  |
|-----------------------------|----------|----------|--------|
| SGDClassifier               | 0.5002   | 0.3390   | 0.5020 |
| MultinomialNB               | 0.5597   | 0.5596   | 0.5596 |
| PassiveAggressiveClassifier | 0.5406   | 0.4661   | 0.5420 |

### 13.3.6 Sprawdźmy metryki dla zbioru uczącego

In [22]:

```
y_true_list=[]
y_pred_list=[]
for i,(batch,labels) in enumerate(ds_train):
    X=batch.numpy()
    X = np.reshape(X,(X.shape[0],-1))
    y=labels.numpy()
    y_pred = clf.predict(X)
    y_true_list.append(y)
    y_pred_list.append(y_pred)

y_true = np.hstack(y_true_list)
# y_true.shape

y_pred=np.hstack(y_pred_list)
y_pred.shape

print(f'Accuracy:{sklearn.metrics.accuracy_score(y_true,y_pred)}')
print(f'F1:{sklearn.metrics.f1_score(y_true,y_pred,average="macro")}')
print(f'auroc:{sklearn.metrics.roc_auc_score(y_true,y_pred,average="macro")}')
```

Accuracy:0.5642080697660137

F1:0.49360421839379365

auroc:0.5645985227193009

In [23]:

```
from sklearn.metrics import confusion_matrix
conf_mat = confusion_matrix(y_true, y_pred)
print(conf_mat)
```

```
[[ 3107 13193]
 [   999 15267]]
```

## 13.4. Sieć neuronowa

Zbudujemy prostą sieć neuronową korzystającą z komponentów platformy TensorFlow/keras. Sieć składa się z:

- warstwy wejściowej (takiej jak rozmiary obrazu, czyli [100,100,3])
- warstwy spłaszczającej do postaci wektora [300000]
- warstwy ukrytej liczącej 100 neuronów z funkcją aktywacji [relu](https://en.wikipedia.org/wiki/Rectifier_(neural_networks)) ([https://en.wikipedia.org/wiki/Rectifier\\_\(neural\\_networks\)](https://en.wikipedia.org/wiki/Rectifier_(neural_networks)))
- warstwy wyjściowej z dwoma neuronami i funkcją aktywacji [softmax](https://en.wikipedia.org/wiki/Softmax_function) ([https://en.wikipedia.org/wiki/Softmax\\_function](https://en.wikipedia.org/wiki/Softmax_function))

Liczba neuronów wyjściowych odpowiada liczbie klas. Odczytujemy z nich prawdopodobieństwo i wybieramy większe z nich.

Alternatywna konfiguracja w przypadku klasyfikacji binarnej to jeden neuron wyjściowy z funkcją aktywacji sigmoid. Przypomina to regresję logistyczną. W takim przypadku wybieramy 0 dla prawdopodobieństwa  $p < 0.5$  oraz 1 dla  $p \geq 0.5$ .

Funkcją kosztu jest *crossentropy* - odpowiednik *logloss* znanej z regresji logistycznej.

**Uwaga** Podobnie jak zwykłe klasyfikatory - to nie jest architektura odpowiednia do klasyfikacji obrazów.

In [24]:

```
from keras.layers import Input, Dense, Flatten
from keras import Model

def model_builder(input_shape, labels_count):
    model_input = Input(input_shape)
    x = Flatten()(model_input)
    x = Dense(input_shape[1], activation = 'relu')(x)
    x = Dense(labels_count, activation='softmax')(x)
    model = Model(model_input, x, name='simple')
    model.compile(optimizer = 'adam', loss = 'sparse_categorical_crossentropy', metrics=['accuracy'])
    return model

model = model_builder((80,80,3),2)
```

Wyświetlmy informacje o modelu. Ma dużo parametrów, ponad 3 000 000.

Wejście ma rozmiar [None, 100, 100, 3] Podczas uczenia to None zostanie zastąpione wymiarami wsadu batch (tu 128).

In [25]:

```

from keras.utils import plot_model

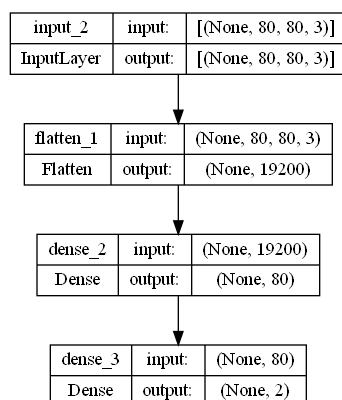
model.summary()
print()
plot_model(model, show_shapes=True, show_layer_names=True, to_file='model.png')
from IPython.display import Image
Image(retina=True, filename='model.png', width = 1000, height = 1000)

```

Model: "simple"

| Layer (type)                | Output Shape        | Param # |
|-----------------------------|---------------------|---------|
| =====                       |                     |         |
| input_2 (InputLayer)        | [(None, 80, 80, 3)] | 0       |
| flatten_1 (Flatten)         | (None, 19200)       | 0       |
| dense_2 (Dense)             | (None, 80)          | 1536080 |
| dense_3 (Dense)             | (None, 2)           | 162     |
| =====                       |                     |         |
| Total params: 1,536,242     |                     |         |
| Trainable params: 1,536,242 |                     |         |
| Non-trainable params: 0     |                     |         |
| =====                       |                     |         |

Out[25]:



Uczymy się.

- Każda epoka to jednokrotna iteracja po zbiorze danych.
- Podział na wsady (batch) jest automatycznie dokonywany przez strumień
- Podczas czytania batcha kolejne obrazy zostaną załadowane z dysku, przeskalowane i umieszczone w pamięci
- Mniej więcej w podobnym czasie poprzedni batch powinien być przetworzony na GPU
- Po każdej epoce następuje walidacja na zbiorze testowym. To przedłuża czas przetwarzania epoki (ale umożliwia przygotowanie wykresów)
- Niestety każda epoka trwa około 37-38 sekund

In [26]:

```
history = model.fit(ds_train, epochs=20, batch_size=128, verbose=1, validation_data=ds_test)
```



Epoch 1/20  
256/256 [=====] - 20s 76ms/step - loss: 1.0  
472 - accuracy: 0.5412 - val\_loss: 0.7275 - val\_accuracy: 0.5641  
Epoch 2/20  
256/256 [=====] - 20s 79ms/step - loss: 0.7  
990 - accuracy: 0.5818 - val\_loss: 0.8316 - val\_accuracy: 0.5350  
Epoch 3/20  
256/256 [=====] - 21s 82ms/step - loss: 0.7  
135 - accuracy: 0.6109 - val\_loss: 0.7920 - val\_accuracy: 0.5528  
Epoch 4/20  
256/256 [=====] - 22s 87ms/step - loss: 0.6  
978 - accuracy: 0.6165 - val\_loss: 0.6954 - val\_accuracy: 0.6025  
Epoch 5/20  
256/256 [=====] - 21s 81ms/step - loss: 0.6  
431 - accuracy: 0.6433 - val\_loss: 0.6824 - val\_accuracy: 0.6075  
Epoch 6/20  
256/256 [=====] - 21s 82ms/step - loss: 0.6  
481 - accuracy: 0.6435 - val\_loss: 0.7245 - val\_accuracy: 0.5938  
Epoch 7/20  
256/256 [=====] - 21s 81ms/step - loss: 0.6  
282 - accuracy: 0.6550 - val\_loss: 0.7575 - val\_accuracy: 0.5777  
Epoch 8/20  
256/256 [=====] - 20s 76ms/step - loss: 0.6  
428 - accuracy: 0.6499 - val\_loss: 0.7445 - val\_accuracy: 0.5797  
Epoch 9/20  
256/256 [=====] - 19s 75ms/step - loss: 0.6  
239 - accuracy: 0.6563 - val\_loss: 0.6400 - val\_accuracy: 0.6340  
Epoch 10/20  
256/256 [=====] - 20s 76ms/step - loss: 0.6  
047 - accuracy: 0.6656 - val\_loss: 0.6826 - val\_accuracy: 0.6114  
Epoch 11/20  
256/256 [=====] - 20s 76ms/step - loss: 0.6  
005 - accuracy: 0.6739 - val\_loss: 0.6663 - val\_accuracy: 0.6161  
Epoch 12/20  
256/256 [=====] - 19s 75ms/step - loss: 0.5  
963 - accuracy: 0.6784 - val\_loss: 0.6886 - val\_accuracy: 0.6164  
Epoch 13/20  
256/256 [=====] - 20s 77ms/step - loss: 0.5  
793 - accuracy: 0.6915 - val\_loss: 0.6927 - val\_accuracy: 0.6146  
Epoch 14/20  
256/256 [=====] - 21s 83ms/step - loss: 0.5  
911 - accuracy: 0.6823 - val\_loss: 0.7020 - val\_accuracy: 0.6080  
Epoch 15/20  
256/256 [=====] - 21s 80ms/step - loss: 0.5  
691 - accuracy: 0.7007 - val\_loss: 0.6725 - val\_accuracy: 0.6207  
Epoch 16/20  
256/256 [=====] - 20s 78ms/step - loss: 0.5  
646 - accuracy: 0.7046 - val\_loss: 0.6459 - val\_accuracy: 0.6414  
Epoch 17/20  
256/256 [=====] - 20s 77ms/step - loss: 0.5  
543 - accuracy: 0.7131 - val\_loss: 0.6516 - val\_accuracy: 0.6371  
Epoch 18/20  
256/256 [=====] - 20s 77ms/step - loss: 0.5  
451 - accuracy: 0.7211 - val\_loss: 0.7432 - val\_accuracy: 0.6025  
Epoch 19/20  
256/256 [=====] - 20s 79ms/step - loss: 0.5  
444 - accuracy: 0.7202 - val\_loss: 0.6549 - val\_accuracy: 0.6383  
Epoch 20/20  
256/256 [=====] - 20s 76ms/step - loss: 0.5  
384 - accuracy: 0.7245 - val\_loss: 0.7738 - val\_accuracy: 0.5991

### **TODO 13.4.1**

- Wyświetl wykresy z danymi zawartymi w history


In [27]:

```
print(history.history.keys())
plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.legend()
plt.grid()
plt.show()


plt.figure()
plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])
plt.legend()
plt.grid()
plt.show()
```

WARNING:matplotlib.legend:No artists with labels found to put in legend. Note that artists whose label start with an underscore are ignored when legend() is called with no argument.

```
dict_keys(['loss', 'accuracy', 'val_loss', 'val_accuracy'])
```



WARNING:matplotlib.legend:No artists with labels found to put in legend. Note that artists whose label start with an underscore are ignored when legend() is called with no argument.



## 13.5. Zbiór danych w pamięci

Spróbujemy przetwarzać zbiór danych w pamięci. To powinno przyspieszyć. Colab ma słabe CPU i kod wykonywany po stronie hosta działa wolno.

**Zresetuj środowisko wykonawcze**

### 13.5.1 Funkcja ładująca zbiór danych do macierzy numpy

In [2]:

```
import matplotlib.pyplot as plt
import numpy as np
import tensorflow as tf
import tensorflow_datasets as tfds
from tensorflow.keras.utils import img_to_array
from tqdm import tqdm

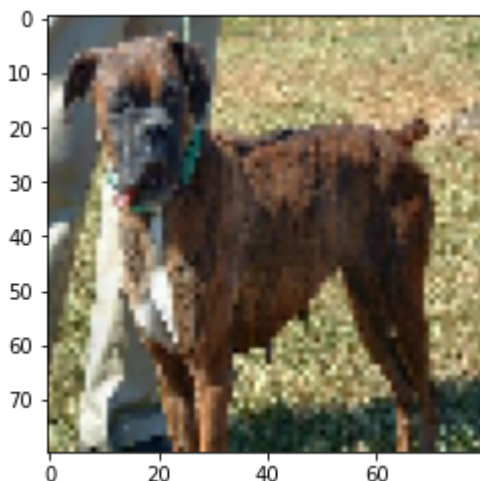
def load_images_to_numpy(name,resize_to=[80,80]):
    ds = tfds.load('cats_vs_dogs', split='train', batch_size=None,as_supervised=True)
    images = []
    labels=[]
    for image, label in tqdm(ds):
        image = tf.image.resize(image, size=resize_to, method=tf.image.ResizeMethod.NEAREST_NEIGHBOR)
        image = img_to_array(image)/255
        images.append(image)
        labels.append(label.numpy())

    # return images,labels
    X=np.array(images)
    y=np.array(labels)
    return X,y
```

In [3]:

```
X,y = load_images_to_numpy('cats_vs_dogs')
plt.imshow(X[0]);
```

100% | ██████████ | 23262/23262 [00:50<00:00, 460.19it/s]



## 13.5.2 Sprawdzamy ile mamy pamięci

In [4]:

```
import psutil  
psutil.virtual_memory()
```

Out[4]:

```
svmem(total=16977661952, available=2503389184, percent=85.3, used=14  
474272768, free=2503389184)
```

## 13.5.3 Duża monolityczna funkcja do ucznia modeli i ich oceny

Napiszemy funkcję, która:

1. Wydzieli zbiór uczący i testowy
2. Wywoła zewnętrzną funkcję do budowy modelu
3. Przeprowadzi uczenie z walidacją
4. Przetestuje model na zbiorze testowym
5. Wypisze metryki

### TODO 13.5.1

- Ad 1. Podziel dane w proporcjach określonych przez test\_size za pomocą train\_test\_split
- Ad 3. Wywołaj model.fit
- Ad 3. Wyświetl historię
- Ad 4. Wyznacz przewidywane etykiety y\_pred
- Ad 5. Oblicz metryki

In [5]:



```

from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score
import matplotlib.pyplot as plt
import psutil

def split_train_test(X, y, model_builder, epochs=20, test_size=0.3, batch_size=128, use_validation_data=True, display_history=True):
    print(psutil.virtual_memory())
    X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=test_size)

    ds_train = tf.data.Dataset.from_tensor_slices((X_train, y_train)).shuffle(1000)
    ds_train = ds_train.batch(batch_size, drop_remainder=True)

    ds_test = tf.data.Dataset.from_tensor_slices((X_test, y_test))
    ds_test = ds_test.batch(batch_size)

    model = model_builder()
    validation_data = None
    if use_validation_data:
        validation_data = ds_test
    history = model.fit(ds_train, epochs=epochs, validation_data=validation_data)

    if display_history:
        print(history.history.keys())
        plt.plot(history.history['loss'])
        plt.plot(history.history['val_loss'])
        plt.legend()
        plt.grid()
        plt.show()

        plt.figure()
        plt.plot(history.history['accuracy'])
        plt.plot(history.history['val_accuracy'])
        plt.legend()
        plt.grid()
        plt.show()

    ds_test = tf.data.Dataset.from_tensor_slices((X_test, y_test))
    ds_test = ds_test.batch(128, drop_remainder=False)

    predict_proba = model.predict(ds_test)
    print(predict_proba.shape)
    y_pred = np.argmax(predict_proba, axis=1)

    acc = accuracy_score(y_test, y_pred)
    prec = precision_score(y_test, y_pred, average='macro')
    recall = recall_score(y_test, y_pred, average='macro')
    f1 = f1_score(y_test, y_pred, average='macro')

    print(f'acc={acc}')
    print(f'prec={prec}')
    print(f'recall={recall}')
    print(f'f1={f1}')

    print(psutil.virtual_memory())
    return model

```

## 13.5.4 Wywołanie

In [32]:


```
m = split_train_test(X,y,lambda: model_builder((80,80,3),2))
```

```
svmem(total=16977661952, available=4780638208, percent=71.8, used=12
197023744, free=4780638208)
Epoch 1/20
127/127 [=====] - 5s 37ms/step - loss: 1.02
14 - accuracy: 0.5487 - val_loss: 0.6759 - val_accuracy: 0.5809
Epoch 2/20
127/127 [=====] - 4s 32ms/step - loss: 0.67
49 - accuracy: 0.5830 - val_loss: 0.6807 - val_accuracy: 0.5731
Epoch 3/20
127/127 [=====] - 4s 31ms/step - loss: 0.66
79 - accuracy: 0.5909 - val_loss: 0.6837 - val_accuracy: 0.5676
Epoch 4/20
127/127 [=====] - 4s 32ms/step - loss: 0.66
48 - accuracy: 0.6009 - val_loss: 0.6603 - val_accuracy: 0.6002
Epoch 5/20
127/127 [=====] - 4s 29ms/step - loss: 0.67
64 - accuracy: 0.5893 - val_loss: 0.6936 - val_accuracy: 0.5584
Epoch 6/20
127/127 [=====] - 4s 30ms/step - loss: 0.65
40 - accuracy: 0.6175 - val_loss: 0.6603 - val_accuracy: 0.5999
Epoch 7/20
127/127 [=====] - 4s 34ms/step - loss: 0.65
80 - accuracy: 0.6059 - val_loss: 0.6712 - val_accuracy: 0.5838
Epoch 8/20
127/127 [=====] - 3s 28ms/step - loss: 0.65
33 - accuracy: 0.6128 - val_loss: 0.6635 - val_accuracy: 0.5949
Epoch 9/20
127/127 [=====] - 4s 32ms/step - loss: 0.66
10 - accuracy: 0.6064 - val_loss: 0.6599 - val_accuracy: 0.6004
Epoch 10/20
127/127 [=====] - 4s 32ms/step - loss: 0.64
31 - accuracy: 0.6291 - val_loss: 0.6556 - val_accuracy: 0.6067
Epoch 11/20
127/127 [=====] - 4s 35ms/step - loss: 0.64
38 - accuracy: 0.6253 - val_loss: 0.6701 - val_accuracy: 0.5838
Epoch 12/20
127/127 [=====] - 4s 34ms/step - loss: 0.64
90 - accuracy: 0.6215 - val_loss: 0.6712 - val_accuracy: 0.5865
Epoch 13/20
127/127 [=====] - 4s 34ms/step - loss: 0.64
42 - accuracy: 0.6243 - val_loss: 0.6574 - val_accuracy: 0.6037
Epoch 14/20
127/127 [=====] - 4s 34ms/step - loss: 0.64
80 - accuracy: 0.6208 - val_loss: 0.7139 - val_accuracy: 0.5488
Epoch 15/20
127/127 [=====] - 4s 32ms/step - loss: 0.64
11 - accuracy: 0.6288 - val_loss: 0.6545 - val_accuracy: 0.6098
Epoch 16/20
127/127 [=====] - 4s 32ms/step - loss: 0.63
69 - accuracy: 0.6352 - val_loss: 0.6546 - val_accuracy: 0.6060
Epoch 17/20
127/127 [=====] - 4s 29ms/step - loss: 0.64
38 - accuracy: 0.6266 - val_loss: 0.6569 - val_accuracy: 0.6075
Epoch 18/20
127/127 [=====] - 3s 27ms/step - loss: 0.62
98 - accuracy: 0.6459 - val_loss: 0.6655 - val_accuracy: 0.5913
Epoch 19/20
127/127 [=====] - 4s 29ms/step - loss: 0.63
40 - accuracy: 0.6414 - val_loss: 0.6614 - val_accuracy: 0.6074
Epoch 20/20
```


```
127/127 [=====] - 4s 33ms/step - loss: 0.6284 - accuracy: 0.6449 - val_loss: 0.6533 - val_accuracy: 0.6123
```

WARNING:matplotlib.legend:No artists with labels found to put in legend. Note that artists whose label start with an underscore are ignored when legend() is called with no argument.

```
dict_keys(['loss', 'accuracy', 'val_loss', 'val_accuracy'])
```



WARNING:matplotlib.legend:No artists with labels found to put in legend. Note that artists whose label start with an underscore are ignored when legend() is called with no argument.



```
55/55 [=====] - 1s 15ms/step
(6979, 2)
acc=0.6122653675311649
prec=0.6125414634516392
recall=0.6120915094409354
f1=0.6118024533612233
svmem(total=16977661952, available=3190104064, percent=81.2, used=13787557888, free=3190104064)
```

### TODO 13.5.2

- Oceń zużycie pamięci
- Czy pamięć została zwolniona po wyjściu z funkcji?

Zużycie pamięci przed: svmem(total=16977661952, available=4780638208, percent=71.8, used=12197023744, free=4780638208)

Zużycie pamięci po:

svmem(total=16977661952, available=3190104064, percent=81.2, used=13787557888, free=3190104064)

Wyraźnie widać że procentowe zużycie pamięci znacznie wzrosło - o około 10 punktów procentowe.

Nie, pamięć nie została zwolniona po wyjściu z funkcji

## 13.6. Konwolucyjna sieć neuronowa CNN

- Sieci konwolucyjne automatycznie ekstrahują cechy z obrazu stosując operacje konwolucji.
- Nie posługują się ustalonymi manualnie wagami filtrów konwolucji, ale wyznaczają je w trakcie uczenia

Zaprezentowane tu modele pochodzą ze strony [How to Classify Photos of Dogs and Cats \(with 97% accuracy\)](https://machinelearningmastery.com/how-to-develop-a-convolutional-neural-network-to-classify-photos-of-dogs-and-cats/) (<https://machinelearningmastery.com/how-to-develop-a-convolutional-neural-network-to-classify-photos-of-dogs-and-cats/>) ale zostały nieco zmodyfikowane

- Rozmiary obrazów są mniejsze (80 x 80) vs (224 x 224). Zapewne wyniki będą gorsze
- Zastosowane są modele z dwoma neuronami wyjściowym

### 13.6.1 Prosty model

In [33]:

```
from keras.layers import Conv2D, Dense, Flatten, MaxPooling2D
from keras import Sequential
from keras.optimizers import SGD
from keras.applications.vgg16 import VGG16
from keras.models import Model

def define_simple_model():
    model = Sequential()
    model.add(Conv2D(32, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same', input_shape=(80, 80, 3)))
    model.add(MaxPooling2D((2, 2)))
    model.add(Flatten())
    model.add(Dense(128, activation='relu', kernel_initializer='he_uniform'))
    model.add(Dense(2, activation='softmax'))
    # compile model
    model.compile(optimizer='adam', loss='sparse_categorical_crossentropy', metrics=['accuracy'])
    return model

model = define_simple_model()
```

#### Inspekcja modelu

In [34]:

```
from keras.utils import plot_model

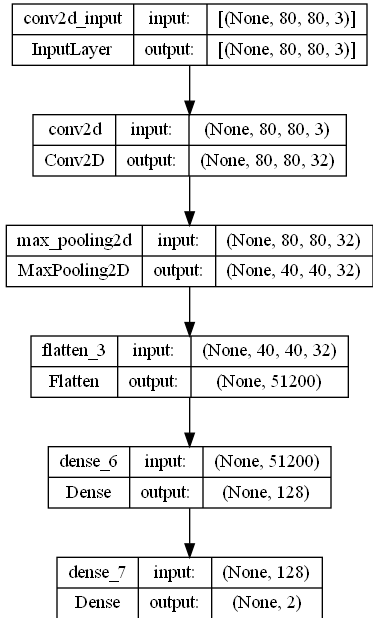
model.summary()
print()
plot_model(model, show_shapes=True, show_layer_names=True, to_file='model.png')
from IPython.display import Image
Image(retina=True, filename='model.png',width = 1000, height = 1000)
```

Model: "sequential"

| Layer (type)                 | Output Shape       | Param # |
|------------------------------|--------------------|---------|
| conv2d (Conv2D)              | (None, 80, 80, 32) | 896     |
| max_pooling2d (MaxPooling2D) | (None, 40, 40, 32) | 0       |
| flatten_3 (Flatten)          | (None, 51200)      | 0       |
| dense_6 (Dense)              | (None, 128)        | 6553728 |
| dense_7 (Dense)              | (None, 2)          | 258     |

=====  
Total params: 6,554,882  
Trainable params: 6,554,882  
Non-trainable params: 0  
=====

Out[34]:



Uczenie i ocena

In [35]:

```
m = split_train_test(X,y,define_simple_model,epochs=20)
```




```
svmem(total=16977661952, available=4925558784, percent=71.0, used=12
052103168, free=4925558784)
Epoch 1/20
127/127 [=====] - 56s 433ms/step - loss: 1.
9912 - accuracy: 0.5418 - val_loss: 0.6775 - val_accuracy: 0.5466
Epoch 2/20
127/127 [=====] - 53s 417ms/step - loss: 0.
6510 - accuracy: 0.6125 - val_loss: 0.6173 - val_accuracy: 0.6550
Epoch 3/20
127/127 [=====] - 52s 411ms/step - loss: 0.
6054 - accuracy: 0.6676 - val_loss: 0.6045 - val_accuracy: 0.6647
Epoch 4/20
127/127 [=====] - 51s 403ms/step - loss: 0.
5717 - accuracy: 0.7029 - val_loss: 0.5625 - val_accuracy: 0.7053
Epoch 5/20
127/127 [=====] - 52s 410ms/step - loss: 0.
5527 - accuracy: 0.7130 - val_loss: 0.5747 - val_accuracy: 0.6965
Epoch 6/20
127/127 [=====] - 51s 402ms/step - loss: 0.
5217 - accuracy: 0.7386 - val_loss: 0.5401 - val_accuracy: 0.7173
Epoch 7/20
127/127 [=====] - 53s 414ms/step - loss: 0.
4961 - accuracy: 0.7608 - val_loss: 0.5223 - val_accuracy: 0.7379
Epoch 8/20
127/127 [=====] - 51s 405ms/step - loss: 0.
4788 - accuracy: 0.7721 - val_loss: 0.5159 - val_accuracy: 0.7438
Epoch 9/20
127/127 [=====] - 53s 416ms/step - loss: 0.
4492 - accuracy: 0.7909 - val_loss: 0.5248 - val_accuracy: 0.7409
Epoch 10/20
127/127 [=====] - 53s 421ms/step - loss: 0.
4391 - accuracy: 0.7950 - val_loss: 0.5146 - val_accuracy: 0.7474
Epoch 11/20
127/127 [=====] - 51s 405ms/step - loss: 0.
4117 - accuracy: 0.8101 - val_loss: 0.5202 - val_accuracy: 0.7461
Epoch 12/20
127/127 [=====] - 51s 404ms/step - loss: 0.
4036 - accuracy: 0.8145 - val_loss: 0.5376 - val_accuracy: 0.7442
Epoch 13/20
127/127 [=====] - 52s 407ms/step - loss: 0.
3855 - accuracy: 0.8250 - val_loss: 0.5166 - val_accuracy: 0.7543
Epoch 14/20
127/127 [=====] - 51s 404ms/step - loss: 0.
3644 - accuracy: 0.8396 - val_loss: 0.5299 - val_accuracy: 0.7527
Epoch 15/20
127/127 [=====] - 51s 406ms/step - loss: 0.
3454 - accuracy: 0.8485 - val_loss: 0.5307 - val_accuracy: 0.7547
Epoch 16/20
127/127 [=====] - 52s 407ms/step - loss: 0.
3237 - accuracy: 0.8612 - val_loss: 0.5512 - val_accuracy: 0.7494
Epoch 17/20
127/127 [=====] - 51s 403ms/step - loss: 0.
3076 - accuracy: 0.8667 - val_loss: 0.5549 - val_accuracy: 0.7534
Epoch 18/20
127/127 [=====] - 51s 402ms/step - loss: 0.
3028 - accuracy: 0.8671 - val_loss: 0.5572 - val_accuracy: 0.7521
Epoch 19/20
127/127 [=====] - 51s 403ms/step - loss: 0.
2809 - accuracy: 0.8800 - val_loss: 0.5903 - val_accuracy: 0.7508
Epoch 20/20
```


```
127/127 [=====] - 51s 404ms/step - loss: 0.2627 - accuracy: 0.8882 - val_loss: 0.6012 - val_accuracy: 0.7511
```

WARNING:matplotlib.legend:No artists with labels found to put in legend. Note that artists whose label start with an underscore are ignored when legend() is called with no argument.

```
dict_keys(['loss', 'accuracy', 'val_loss', 'val_accuracy'])
```



WARNING:matplotlib.legend:No artists with labels found to put in legend. Note that artists whose label start with an underscore are ignored when legend() is called with no argument.



55/55 [=====] - 6s 98ms/step

(6979, 2)

acc=0.7511104742799828

prec=0.7538574243172546

recall=0.7518064901397714

f1=0.7507429121935926

svmem(total=16977661952, available=4906016768, percent=71.1, used=12071645184, free=4906016768)

## 13.6.2 Zaawansowany model korzystający z transfer learning

In [7]:

```
from keras.layers import Conv2D, Dense, Flatten, MaxPooling2D
from keras import Sequential
from keras.optimizers import SGD
from keras.applications.vgg16 import VGG16
from keras.models import Model

# Na podstawie
# https://machinelearningmastery.com/how-to-develop-a-convolutional-neural-netwo
rk-to-classify-photos-of-dogs-and-cats/

def define_vgg16_model():
    # load model
    model = VGG16(include_top=False, input_shape=(80, 80, 3))
    # mark loaded layers as not trainable
    for layer in model.layers:
        layer.trainable = False
    # add new classifier layers
    flat1 = Flatten()(model.layers[-1].output)
    class1 = Dense(128, activation='relu', kernel_initializer='he_uniform')(flat1)
    output = Dense(2, activation='softmax')(class1)
    # define new model
    model = Model(inputs=model.inputs, outputs=output)
    # compile model
    model.compile(optimizer='adam', loss='sparse_categorical_crossentropy',
metrics=['accuracy'])
    return model

model = define_vgg16_model()
```

Inspekcja modelu

In [37]:

```
from keras.utils import plot_model

model.summary()
print()
plot_model(model, show_shapes=True, show_layer_names=True, to_file='model.png')
from IPython.display import Image
Image(retina=True, filename='model.png', width = 1000, height = 1000)
```

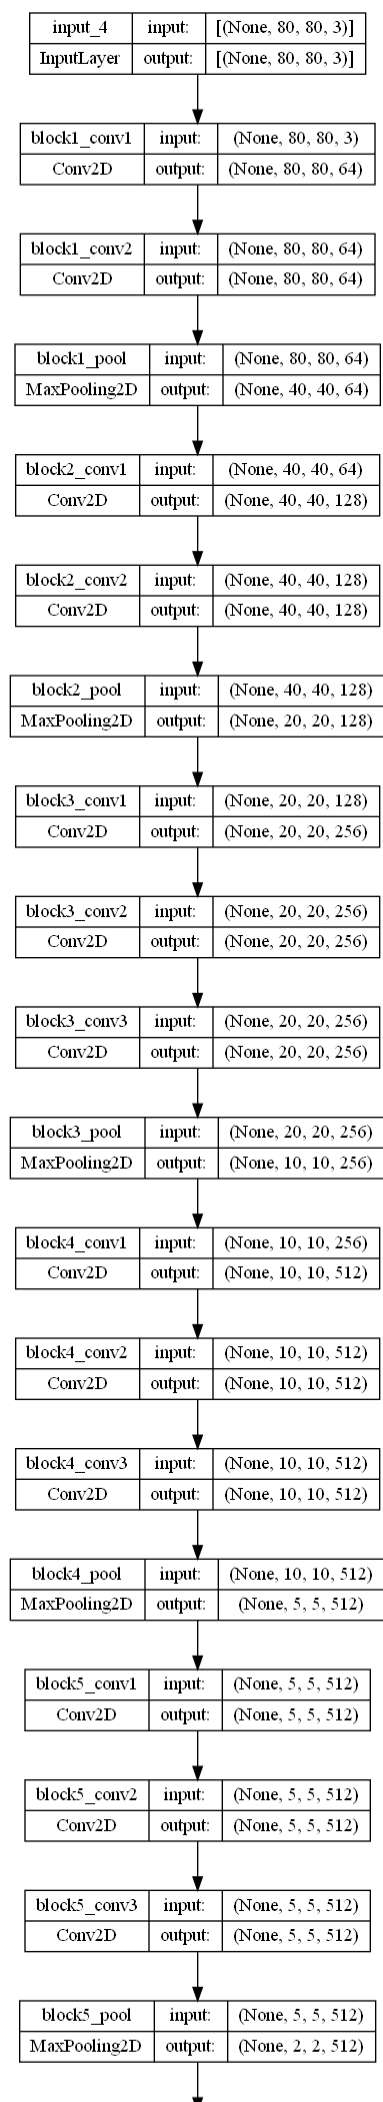
Model: "model"

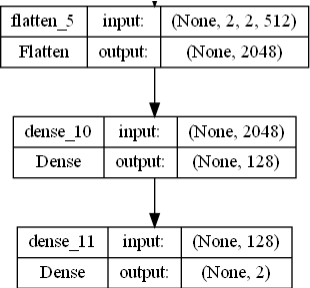
| Layer (type)                     | Output Shape        | Param # |
|----------------------------------|---------------------|---------|
| input_4 (InputLayer)             | [(None, 80, 80, 3)] | 0       |
| block1_conv1 (Conv2D)            | (None, 80, 80, 64)  | 1792    |
| block1_conv2 (Conv2D)            | (None, 80, 80, 64)  | 36928   |
| block1_pool (MaxPooling2D)       | (None, 40, 40, 64)  | 0       |
| block2_conv1 (Conv2D)            | (None, 40, 40, 128) | 73856   |
| block2_conv2 (Conv2D)            | (None, 40, 40, 128) | 147584  |
| block2_pool (MaxPooling2D)       | (None, 20, 20, 128) | 0       |
| block3_conv1 (Conv2D)            | (None, 20, 20, 256) | 295168  |
| block3_conv2 (Conv2D)            | (None, 20, 20, 256) | 590080  |
| block3_conv3 (Conv2D)            | (None, 20, 20, 256) | 590080  |
| block3_pool (MaxPooling2D)       | (None, 10, 10, 256) | 0       |
| block4_conv1 (Conv2D)            | (None, 10, 10, 512) | 1180160 |
| block4_conv2 (Conv2D)            | (None, 10, 10, 512) | 2359808 |
| block4_conv3 (Conv2D)            | (None, 10, 10, 512) | 2359808 |
| block4_pool (MaxPooling2D)       | (None, 5, 5, 512)   | 0       |
| block5_conv1 (Conv2D)            | (None, 5, 5, 512)   | 2359808 |
| block5_conv2 (Conv2D)            | (None, 5, 5, 512)   | 2359808 |
| block5_conv3 (Conv2D)            | (None, 5, 5, 512)   | 2359808 |
| block5_pool (MaxPooling2D)       | (None, 2, 2, 512)   | 0       |
| flatten_5 (Flatten)              | (None, 2048)        | 0       |
| dense_10 (Dense)                 | (None, 128)         | 262272  |
| dense_11 (Dense)                 | (None, 2)           | 258     |
| Total params: 14,977,218         |                     |         |
| Trainable params: 262,530        |                     |         |
| Non-trainable params: 14,714,688 |                     |         |

Out[37]:









In [38]:

```
m = split_train_test(X,y,define_vgg16_model,epochs=20)
```

svmm(total=16977661952, available=6163402752, percent=63.7, used=10814259200, free=6163402752)

Epoch 1/20

127/127 [=====] - 348s 3s/step - loss: 0.44  
32 - accuracy: 0.7849 - val\_loss: 0.3776 - val\_accuracy: 0.8253

Epoch 2/20

127/127 [=====] - 341s 3s/step - loss: 0.35  
46 - accuracy: 0.8401 - val\_loss: 0.3559 - val\_accuracy: 0.8356

Epoch 3/20

127/127 [=====] - 344s 3s/step - loss: 0.33  
41 - accuracy: 0.8484 - val\_loss: 0.3477 - val\_accuracy: 0.8431

Epoch 4/20

127/127 [=====] - 345s 3s/step - loss: 0.30  
63 - accuracy: 0.8640 - val\_loss: 0.3430 - val\_accuracy: 0.8470

Epoch 5/20

127/127 [=====] - 358s 3s/step - loss: 0.29  
09 - accuracy: 0.8734 - val\_loss: 0.3439 - val\_accuracy: 0.8481

Epoch 6/20

127/127 [=====] - 347s 3s/step - loss: 0.26  
31 - accuracy: 0.8882 - val\_loss: 0.3580 - val\_accuracy: 0.8404

Epoch 7/20

127/127 [=====] - 366s 3s/step - loss: 0.24  
20 - accuracy: 0.8982 - val\_loss: 0.3656 - val\_accuracy: 0.8421

Epoch 8/20

127/127 [=====] - 378s 3s/step - loss: 0.21  
99 - accuracy: 0.9110 - val\_loss: 0.3584 - val\_accuracy: 0.8454

Epoch 9/20

127/127 [=====] - 338s 3s/step - loss: 0.19  
94 - accuracy: 0.9196 - val\_loss: 0.3710 - val\_accuracy: 0.8431

Epoch 10/20

127/127 [=====] - 331s 3s/step - loss: 0.17  
44 - accuracy: 0.9320 - val\_loss: 0.3934 - val\_accuracy: 0.8421

Epoch 11/20

127/127 [=====] - 338s 3s/step - loss: 0.15  
06 - accuracy: 0.9461 - val\_loss: 0.3965 - val\_accuracy: 0.8460

Epoch 12/20

127/127 [=====] - 339s 3s/step - loss: 0.13  
65 - accuracy: 0.9524 - val\_loss: 0.4080 - val\_accuracy: 0.8420

Epoch 13/20

127/127 [=====] - 343s 3s/step - loss: 0.11  
75 - accuracy: 0.9623 - val\_loss: 0.4224 - val\_accuracy: 0.8457

Epoch 14/20

127/127 [=====] - 339s 3s/step - loss: 0.10  
52 - accuracy: 0.9658 - val\_loss: 0.4401 - val\_accuracy: 0.8442

Epoch 15/20

127/127 [=====] - 340s 3s/step - loss: 0.09  
29 - accuracy: 0.9699 - val\_loss: 0.4447 - val\_accuracy: 0.8454

Epoch 16/20

127/127 [=====] - 340s 3s/step - loss: 0.07  
51 - accuracy: 0.9791 - val\_loss: 0.4807 - val\_accuracy: 0.8458

Epoch 17/20

127/127 [=====] - 339s 3s/step - loss: 0.06  
75 - accuracy: 0.9816 - val\_loss: 0.4840 - val\_accuracy: 0.8401

Epoch 18/20

127/127 [=====] - 341s 3s/step - loss: 0.05  
63 - accuracy: 0.9867 - val\_loss: 0.5204 - val\_accuracy: 0.8412

Epoch 19/20


127/127 [=====] - 338s 3s/step - loss: 0.04  
75 - accuracy: 0.9911 - val\_loss: 0.5211 - val\_accuracy: 0.8379

Epoch 20/20


```
127/127 [=====] - 329s 3s/step - loss: 0.0418 - accuracy: 0.9914 - val_loss: 0.5296 - val_accuracy: 0.8421
```

WARNING:matplotlib.legend:No artists with labels found to put in legend. Note that artists whose label start with an underscore are ignored when legend() is called with no argument.

```
dict_keys(['loss', 'accuracy', 'val_loss', 'val_accuracy'])
```



WARNING:matplotlib.legend:No artists with labels found to put in legend. Note that artists whose label start with an underscore are ignored when legend() is called with no argument.



```
55/55 [=====] - 118s 2s/step
(6979, 2)
acc=0.8420977217366384
prec=0.8428989329964571
recall=0.8421772289162761
f1=0.8420257150979455
svmmem(total=16977661952, available=2992447488, percent=82.4, used=13
985214464, free=2992447488)
```

### 13.6.3 Zobaczmy koty i psy...

In [39]:

```
ds = tfds.load('cats_vs_dogs', split='train', batch_size=None, as_supervised=True)
ds.shuffle(1000)
it = ds.as_numpy_iterator()
plt.rcParams["figure.figsize"] = (10,10)
labels=['Cat','Dog']
input_size=80
for i in range(9):
    ax = plt.subplot(330 + 1 + i)
    image, label = next(it)
    image = tf.image.resize(image, size=[80,80], method=tf.image.ResizeMethod.NEAREST_NEIGHBOR)
    batched_image = tf.expand_dims(image, axis=0)
    pp = m.predict(batched_image)
    c = np.argmax(pp)
    ax.set_title(labels[label]+" - predicted:"+labels[c])
    ax.imshow(image)
plt.show()
```

```

1/1 [=====] - 1s 651ms/step
1/1 [=====] - 0s 112ms/step
1/1 [=====] - 0s 105ms/step
1/1 [=====] - 0s 92ms/step
1/1 [=====] - 0s 101ms/step
1/1 [=====] - 0s 69ms/step
1/1 [=====] - 0s 61ms/step
1/1 [=====] - 0s 52ms/step
1/1 [=====] - 0s 65ms/step

```



### TODO 13.6.1

- Przeanalizuj wykresy. Ile epok wystarczyło wybrać?
- Jak zmienia się wartość funkcji kosztu i trafność dla zbioru uczącego?
- Jak scharakteryzujesz sytuację począwszy od 5-6 epoki?



Wystarczyło wybrać 5 epok.

Funkcja kosztu maleje przez cały okres trenowania sieci. Trafność dla zbioru uczącego rośnie logarytmicznie.

Od 5-6 epoki `val_loss` dla zbioru testowego zaczyna rosnąć. Parametr `val_loss` odnosi się do wartości funkcji straty obliczanej na zestawie walidacyjnym podczas procesu uczenia modelu maszynowego. Celem uczenia maszynowego jest minimalizacja funkcji straty, a więc uzyskanie malejącego parametru `val_loss`.

## 13.6.4 Twój model

### TODO 13.6.2

- Wybierz jeden z modeli ze strony [How to Classify Photos of Dogs and Cats \(with 97% accuracy\)](https://machinelearningmastery.com/how-to-develop-a-convolutional-neural-network-to-classify-photos-of-dogs-and-cats/) (<https://machinelearningmastery.com/how-to-develop-a-convolutional-neural-network-to-classify-photos-of-dogs-and-cats/>)
- Dostosuj analogicznie do innej reprezentacji wyjścia (2 neurony)
- Wyświetl strukturę
- Przeprowadź uczenie i ocenę

In [8]:

```
def define_model():
    model = Sequential()
    model.add(Conv2D(32, (3, 3), activation='relu', kernel_initializer='he_uniform', padding='same', input_shape=(80, 80, 3)))
    model.add(MaxPooling2D((2, 2)))
    model.add(Flatten())
    model.add(Dense(128, activation='relu', kernel_initializer='he_uniform'))
    model.add(Dense(1, activation='sigmoid'))
    opt = SGD(lr=0.001, momentum=0.9)
    model.compile(optimizer=opt, loss='binary_crossentropy', metrics=['accuracy'])
    return model

model = define_model()
```

```
c:\Users\krzyc\anaconda3-1\lib\site-packages\keras\optimizers\optimizer_v2\gradient_descent.py:111: UserWarning: The `lr` argument is deprecated, use `learning_rate` instead.
  super().__init__(name, **kwargs)
```

# Inspekcja modelu

In [10]:

```
from keras.utils import plot_model

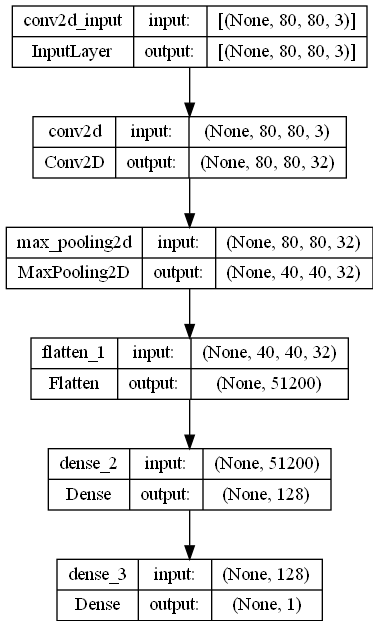
model.summary()
print()
plot_model(model, show_shapes=True, show_layer_names=True, to_file='model.png')
from IPython.display import Image
Image(retina=True, filename='model.png',width = 1000, height = 1000)
```

Model: "sequential"

| Layer (type)                 | Output Shape       | Param # |
|------------------------------|--------------------|---------|
| conv2d (Conv2D)              | (None, 80, 80, 32) | 896     |
| max_pooling2d (MaxPooling2D) | (None, 40, 40, 32) | 0       |
| flatten_1 (Flatten)          | (None, 51200)      | 0       |
| dense_2 (Dense)              | (None, 128)        | 6553728 |
| dense_3 (Dense)              | (None, 1)          | 129     |

=====  
Total params: 6,554,753  
Trainable params: 6,554,753  
Non-trainable params: 0  
=====

Out[10]:



## Uczenie i ocena

### TODO 13.6.3

- Porównaj złożoność wybranego modelu z innymi i skomentuj wyniki

In [12]:


```
m = split_train_test(X,y,define_vgg16_model,epochs=20)
```

```
svmem(total=16977661952, available=874242048, percent=94.9, used=161
03419904, free=874242048)
Epoch 1/20
127/127 [=====] - 343s 3s/step - loss: 0.41
10 - accuracy: 0.8087 - val_loss: 0.3809 - val_accuracy: 0.8281
Epoch 2/20
127/127 [=====] - 347s 3s/step - loss: 0.34
03 - accuracy: 0.8479 - val_loss: 0.3628 - val_accuracy: 0.8374
Epoch 3/20
127/127 [=====] - 340s 3s/step - loss: 0.32
83 - accuracy: 0.8537 - val_loss: 0.3680 - val_accuracy: 0.8339
Epoch 4/20
127/127 [=====] - 333s 3s/step - loss: 0.29
94 - accuracy: 0.8669 - val_loss: 0.3584 - val_accuracy: 0.8399
Epoch 5/20
127/127 [=====] - 334s 3s/step - loss: 0.27
71 - accuracy: 0.8815 - val_loss: 0.3615 - val_accuracy: 0.8407
Epoch 6/20
127/127 [=====] - 333s 3s/step - loss: 0.25
71 - accuracy: 0.8920 - val_loss: 0.3684 - val_accuracy: 0.8369
Epoch 7/20
127/127 [=====] - 333s 3s/step - loss: 0.24
13 - accuracy: 0.8970 - val_loss: 0.3679 - val_accuracy: 0.8394
Epoch 8/20
127/127 [=====] - 334s 3s/step - loss: 0.20
41 - accuracy: 0.9173 - val_loss: 0.3893 - val_accuracy: 0.8325
Epoch 9/20
127/127 [=====] - 334s 3s/step - loss: 0.18
50 - accuracy: 0.9278 - val_loss: 0.3987 - val_accuracy: 0.8389
Epoch 10/20
127/127 [=====] - 337s 3s/step - loss: 0.17
00 - accuracy: 0.9330 - val_loss: 0.4062 - val_accuracy: 0.8308
Epoch 11/20
127/127 [=====] - 338s 3s/step - loss: 0.14
46 - accuracy: 0.9472 - val_loss: 0.4289 - val_accuracy: 0.8269
Epoch 12/20
127/127 [=====] - 335s 3s/step - loss: 0.11
85 - accuracy: 0.9600 - val_loss: 0.4347 - val_accuracy: 0.8319
Epoch 13/20
127/127 [=====] - 331s 3s/step - loss: 0.10
94 - accuracy: 0.9634 - val_loss: 0.4534 - val_accuracy: 0.8345
Epoch 14/20
127/127 [=====] - 333s 3s/step - loss: 0.08
71 - accuracy: 0.9733 - val_loss: 0.4807 - val_accuracy: 0.8319
Epoch 15/20
127/127 [=====] - 333s 3s/step - loss: 0.07
36 - accuracy: 0.9803 - val_loss: 0.4994 - val_accuracy: 0.8286
Epoch 16/20
127/127 [=====] - 341s 3s/step - loss: 0.06
16 - accuracy: 0.9852 - val_loss: 0.5043 - val_accuracy: 0.8335
Epoch 17/20
127/127 [=====] - 341s 3s/step - loss: 0.05
27 - accuracy: 0.9880 - val_loss: 0.5334 - val_accuracy: 0.8356
Epoch 18/20
127/127 [=====] - 339s 3s/step - loss: 0.04
15 - accuracy: 0.9929 - val_loss: 0.5639 - val_accuracy: 0.8288
Epoch 19/20
127/127 [=====] - 338s 3s/step - loss: 0.03
81 - accuracy: 0.9931 - val_loss: 0.5751 - val_accuracy: 0.8331
Epoch 20/20
```


127/127 [=====] - 331s 3s/step - loss: 0.0300 - accuracy: 0.9963 - val\_loss: 0.6036 - val\_accuracy: 0.8315

WARNING:matplotlib.legend:No artists with labels found to put in legend. Note that artists whose label start with an underscore are ignored when legend() is called with no argument.

dict\_keys(['loss', 'accuracy', 'val\_loss', 'val\_accuracy'])



WARNING:matplotlib.legend:No artists with labels found to put in legend. Note that artists whose label start with an underscore are ignored when legend() is called with no argument.



55/55 [=====] - 101s 2s/step

(6979, 2)

acc=0.831494483450351

prec=0.8318805726689342

recall=0.8316826730692277

f1=0.8314840174660438

svmmem(total=16977661952, available=2973536256, percent=82.5, used=14004125696, free=2973536256)

In [ ]:

```
ds = tfds.load('cats_vs_dogs', split='train', batch_size=None, as_supervised=True).shuffle(1000)
it = ds.as_numpy_iterator()
plt.rcParams["figure.figsize"] = (10,10)
labels=['Cat', 'Dog']
input_size=100
for i in range(9):
    ax = plt.subplot(330 + 1 + i)
    image, label = next(it)
    image = tf.image.resize(image, size=[100,100], method=tf.image.ResizeMethod.NEAREST_NEIGHBOR)
    batched_image = tf.expand_dims(image, axis=0)
    pp = m.predict(batched_image)
    c = np.argmax(pp)
    ax.set_title(labels[label]+" - predicted:"+labels[c])
    ax.imshow(image)
plt.show()
```

Porównanie trzech ostatnich modeli w tym notatniku

Model 1:

Liczba warstw konwolucyjnych: 1 Liczba warstw poolingowych: 1 Liczba warstw gęstych: 2 Całkowita liczba warstw: 6

Model 2:

Liczba warstw konwolucyjnych: 1 Liczba warstw poolingowych: 1 Liczba warstw gęstych: 1 Liczba warstw wyjściowych: 1 Całkowita liczba warstw: 5

Model 3:

Liczba warstw konwolucyjnych: 1 Liczba warstw poolingowych: 1 Liczba warstw gęstych: 2 Liczba warstw wyjściowych: 1 Całkowita liczba warstw: 6